

AHDB1_Death_TS

Tonia S. Schwartz

2025-01-22

https://github.com/StatQuest/logistic_regression_demo/blob/master/logistic_regression_demo.R
###Read in main data files

```
library(multcomp)
library(ggplot2)
library(nlme)
library(grid)
library(gridExtra)
library(emmeans)
library(cowplot)
library(ggplot2)
```

Read in the data

```
data=read.csv("AHDB_MasterDataSheet.csv")
data
```

##	AgeCategory	ID_Band	Sex	Treatment	Died_InTrt	Treatment_Category
## 1	Younger	4741	F	B	0	Heat
## 2	Younger	4703	F	B	0	Heat
## 3	Younger	4712	F	B	0	Heat
## 4	Younger	4787	M	B	0	Heat
## 5	Younger	4791	M	B	0	Heat
## 6	Younger	4783	M	B	0	Heat
## 7	Younger	4757	M	B	0	Heat
## 8	Older	4593	M	B	0	Heat
## 9	Older	4534	F	B	0	Heat
## 10	Younger	4792	F	B	0	Heat
## 11	Younger	4769	M	B	0	Heat
## 12	Older	4600	M	B	0	Heat
## 13	Older	4614	M	B	0	Heat
## 14	Younger	4816	M	B	0	Heat
## 15	Younger	4778	M	B	0	Heat
## 16	Younger	4797	F	B	0	Heat
## 17	Older	4613	M	B	0	Heat
## 18	Younger	4788	M	B	0	Heat
## 19	Older	4620	F	B	0	Heat
## 20	Younger	4771	M	B	0	Heat

## 21	Younger	4764	M	B	0	Heat
## 22	Younger	4725	M	B	0	Heat
## 23	Younger	4800	M	B	0	Heat
## 24	Younger	4803	M	B	0	Heat
## 25	Younger	4809	F	B	0	Heat
## 26	Younger	4804	M	B	0	Heat
## 27	Older	4531	F	B	1	Heat
## 28	Younger	4768	F	B	1	Heat
## 29	Younger	4761	F	B	1	Heat
## 30	Younger	4774	F	B	1	Heat
## 31	Younger	4802	F	B	1	Heat
## 32	Younger	4810	F	B	1	Heat
## 33	Younger	4817	F	B	1	Heat
## 34	Older	4529	F	C	0	Heat
## 35	Younger	4704	F	C	0	Heat
## 36	Older	4715	F	C	0	Heat
## 37	Older	4608	F	C	0	Heat
## 38	Older	4559	M	C	0	Heat
## 39	Younger	4782	F	C	0	Heat
## 40	Younger	4748	F	C	0	Heat
## 41	Younger	4717	F	C	0	Heat
## 42	Older	4561	F	C	0	Heat
## 43	Younger	4731	F	C	0	Heat
## 44	Older	4554	F	C	0	Heat
## 45	Older	4601	F	C	0	Heat
## 46	Younger	4814	F	C	0	Heat
## 47	Older	4586	F	C	0	Heat
## 48	Younger	4776	F	C	0	Heat
## 49	Younger	4805	F	C	0	Heat
## 50	Younger	4775	F	C	1	Heat
## 51	Younger	4737	F	C	1	Heat
## 52	Older	4530	F	D	0	Heat
## 53	Younger	4759	F	D	0	Heat
## 54	Older	4551	F	D	0	Heat
## 55	Younger	4749	F	D	0	Heat
## 56	Older	4573	F	D	0	Heat
## 57	Older	4578	F	D	0	Heat
## 58	Older	4604	F	D	0	Heat
## 59	Younger	4723	F	D	0	Heat
## 60	Younger	4735	F	D	0	Heat
## 61	Younger	4719	F	D	1	Heat
## 62	Older	4562	F	D	0	Heat
## 63	Younger	4729	F	D	0	Heat
## 64	Younger	4795	F	D	0	Heat
## 65	Older	4621	F	D	0	Heat
## 66	Younger	4773	F	D	0	Heat
## 67	Younger	4752	F	D	0	Heat
## 68	Younger	4765	F	D	0	Heat
## 69	Younger	4798	F	D	1	Heat
## 70	Younger	4739	F	E	0	Heat
## 71	Younger	4745	F	E	0	Heat
## 72	Younger	4785	F	E	0	Heat
## 73	Younger	4784	F	E	0	Heat
## 74	Younger	4758	F	E	0	Heat

```
## 75      Older    4533  F      E      0      Heat
## 76      Older    4579  F      E      0      Heat
## 77      Younger  4738  F      E      0      Heat
## 78      Younger  4582  F      E      0      Heat
## 79      Older    4535  F      E      0      Heat
## 80      Younger  4777  F      E      0      Heat
## 81      Younger  4720  F      E      0      Heat
## 82      Older    4612  F      E      0      Heat
## 83      Older    4548  F      E      0      Heat
## 84      Younger  4625  F      E      0      Heat
## 85      Older    4619  F      E      0      Heat
## 86      Younger  4772  F      E      0      Heat
## 87      Younger  4766  F      E      0      Heat
```

```
xtabs(~ Died_InTrt + Sex, data=data) # Summarize the number of deaths by Sex
```

```
##           Sex
## Died_InTrt F  M
##           0 57 19
##           1 11  0
```

```
xtabs(~ Died_InTrt + AgeCategory, data=data) # Summarize the number of deaths by Age
```

```
##           AgeCategory
## Died_InTrt Older Younger
##           0    27     49
##           1     1     10
```

Run generalized linear model for Sex

```
logistic <- glm(Died_InTrt ~ Sex, data=data, family="binomial")
summary(logistic)
```

```
##
## Call:
## glm(formula = Died_InTrt ~ Sex, family = "binomial", data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.59406  -0.59406  -0.59406  -0.00013   1.90872
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -1.6452     0.3293  -4.996 5.87e-07 ***
## SexM          -16.9209    1496.3960  -0.011   0.991
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
```

```
##
## Null deviance: 66.043 on 86 degrees of freedom
## Residual deviance: 60.192 on 85 degrees of freedom
## AIC: 64.192
##
## Number of Fisher Scoring iterations: 17
```

```
female.log.odds <- log(9 / 36)
female.log.odds
```

```
## [1] -1.386294
```

```
male.log.odds <- log(0 / 18)
male.log.odds
```

```
## [1] -Inf
```

Now calculate the overall “Pseudo R-squared” and its p-value for Sex

```
## NOTE: Since we are doing logistic regression...
## Null deviance = 2*(0 - LogLikelihood(null model))
##              = -2*LogLikelihood(null model)
## Residual deviance = 2*(0 - LogLikelihood(proposed model))
##              = -2*LogLikelihood(proposed model)
ll.null <- logistic$null.deviance/-2
ll.proposed <- logistic$deviance/-2

## McFadden's Pseudo R^2 = [ LL(Null) - LL(Proposed) ] / LL(Null)
(ll.null - ll.proposed) / ll.null
```

```
## [1] 0.08859916
```

```
## chi-square value = 2*(LL(Proposed) - LL(Null))
## p-value = 1 - pchisq(chi-square value, df = 2-1)
1 - pchisq(2*(ll.proposed - ll.null), df=1)
```

```
## [1] 0.01556518
```

```
1 - pchisq((logistic$null.deviance - logistic$deviance), df=1)
```

```
## [1] 0.01556518
```


Lastly, let’s see what this logistic regression predicts, given that an individual is either female or male (and no other data about them).

```

predicted.data_sex <- data.frame(
  probability.of.Data=logistic$fitted.values,
  sex=data$Sex)

## We can plot the data...
ggplot(data=predicted.data_sex, aes(x=sex, y=probability.of.Data)) +
  geom_point(aes(color=sex), size=5) +
  xlab("Sex") +
  ylab("Predicted probability of Death in Acute Heat in all birds") +
  theme_classic()

```



```

## Since there are only two probabilities (one for females and one for males),
## we can use a table to summarize the predicted probabilities.
xtabs(~ probability.of.Data + sex, data=predicted.data_sex)

```

```

##           sex
## probability.of.Data    F  M
## 8.64686924501238e-09  0 19
## 0.161764705882353    68  0

```

Run a generalized linear model for Age

```
logistic2 <- glm(Died_InTrt ~ AgeCategory, data=data, family="binomial")
summary(logistic2)
```

```
##
## Call:
## glm(formula = Died_InTrt ~ AgeCategory, family = "binomial",
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6095  -0.6095  -0.6095  -0.2697   2.5816
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      -3.296      1.018  -3.236  0.00121 **
## AgeCategoryYounger  1.707      1.076   1.586  0.11267
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 66.043  on 86  degrees of freedom
## Residual deviance: 62.328  on 85  degrees of freedom
## AIC: 66.328
##
## Number of Fisher Scoring iterations: 6
```

Now calculate the overall “Pseudo R-squared” and its p-value for Age

```
## NOTE: Since we are doing logistic regression...
## Null deviance = 2*(0 - LogLikelihood(null model))
##              = -2*LogLikelihood(null model)
## Residual deviance = 2*(0 - LogLikelihood(proposed model))
##              = -2*LogLikelihood(proposed model)
ll.null <- logistic2$null.deviance/-2
ll.proposed <- logistic2$deviance/-2

## McFadden's Pseudo R^2 = [ LL(Null) - LL(Proposed) ] / LL(Null)
(ll.null - ll.proposed) / ll.null
```

```
## [1] 0.05625529
```

```
## chi-square value = 2*(LL(Proposed) - LL(Null))
## p-value = 1 - pchisq(chi-square value, df = 2-1)
1 - pchisq(2*(ll.proposed - ll.null), df=1)
```

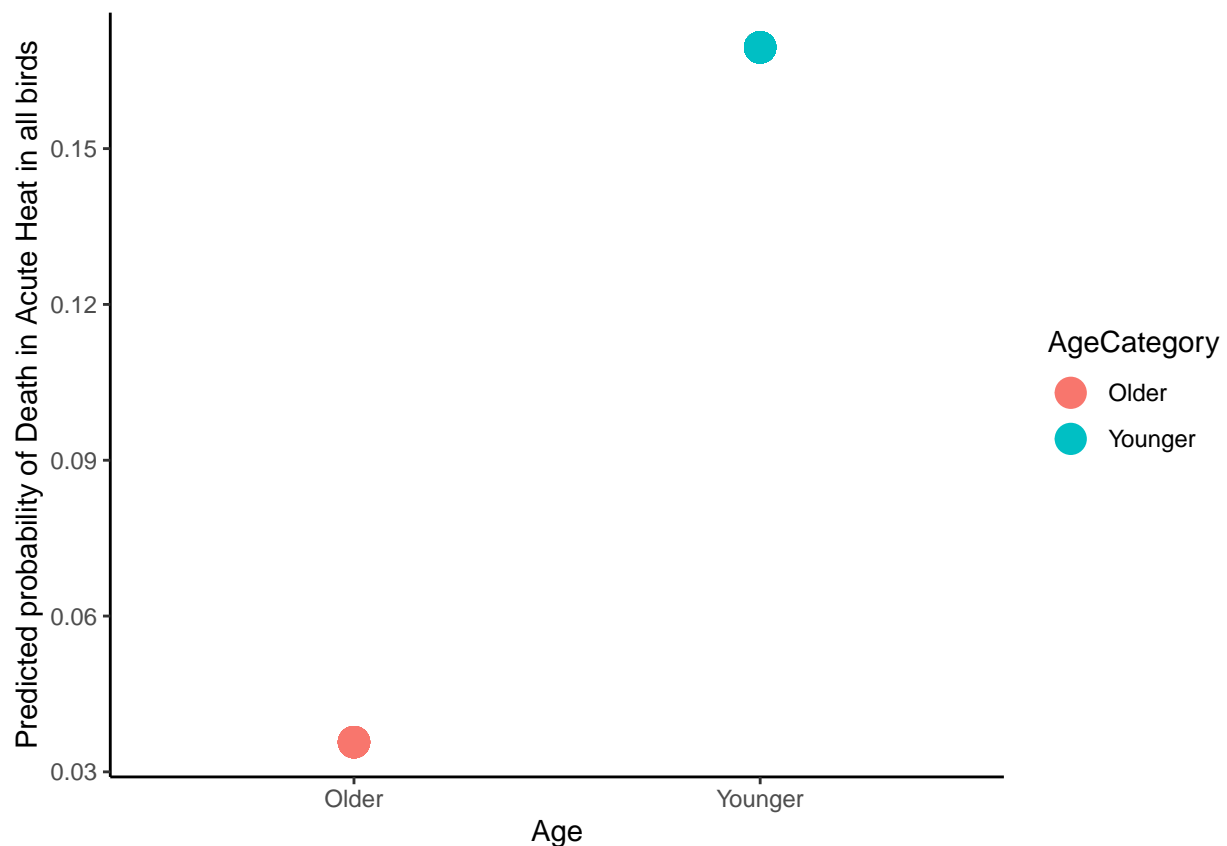
```
## [1] 0.05391723
```

```
1 - pchisq((logistic2$null.deviance - logistic2$deviance), df=1)
```

```
## [1] 0.05391723
```

Lastly, let's see what this logistic regression predicts, given that an individual is either younger or older (and no other data about them).

```
predicted.data_Age <- data.frame(  
  probability.of.Data=logistic2$fitted.values,  
  AgeCategory=data$AgeCategory)  
  
## We can plot the data...  
ggplot(data=predicted.data_Age, aes(x=AgeCategory, y=probability.of.Data)) +  
  geom_point(aes(color=AgeCategory), size=5) +  
  xlab("Age") +  
  ylab("Predicted probability of Death in Acute Heat in all birds") +  
  theme_classic()
```



```
## Since there are only two probabilities (one for females and one for males),
## we can use a table to summarize the predicted probabilities.
xtabs(~ probability.of.Data + AgeCategory, data=predicted.data_Age)
```

```
##
##           AgeCategory
## probability.of.Data Older Younger
##    0.0357142857142918    28      0
##    0.169491525423729     0     59
```

Run a generalized linear model for Sex and Age

```
logistic_Age_Sex <- glm(Died_InTrt ~ AgeCategory + Sex, data=data, family="binomial")
summary(logistic_Age_Sex)
```

```
##
## Call:
## glm(formula = Died_InTrt ~ AgeCategory + Sex, family = "binomial",
##      data = data)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.70896 -0.70896 -0.29817 -0.00009  2.50419
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    -3.091      1.022  -3.023  0.0025 **
## AgeCategoryYounger  1.838      1.084   1.697  0.0898 .
## SexM           -17.999    2388.807  -0.008  0.9940
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 66.043  on 86  degrees of freedom
## Residual deviance: 55.900  on 84  degrees of freedom
## AIC: 61.9
##
## Number of Fisher Scoring iterations: 18
```

Now calculate the overall “Pseudo R-squared” and its p-value for sex and age

```
## NOTE: Since we are doing logistic regression...
## Null deviance = 2*(0 - LogLikelihood(null model))
##              = -2*LogLikelihood(null model)
## Residual deviance = 2*(0 - LogLikelihood(proposed model))
##              = -2*LogLikelihood(proposed model)
```



```

ll.null <- logistic_Age_Sex$null.deviance/-2
ll.proposed <- logistic_Age_Sex$deviance/-2

## McFadden's Pseudo R^2 = [ LL(Null) - LL(Proposed) ] / LL(Null)
(ll.null - ll.proposed) / ll.null

```

```
## [1] 0.1535734
```

```

## chi-square value = 2*(LL(Proposed) - LL(Null))
## p-value = 1 - pchisq(chi-square value, df = 2-1)
1 - pchisq(2*(ll.proposed - ll.null), df=1)

```

```
## [1] 0.001448954
```

```
1 - pchisq((logistic_Age_Sex$null.deviance - logistic_Age_Sex$deviance), df=1)
```

```
## [1] 0.001448954
```


Lastly, let's see what this logistic regression predicts, given that an individual is either younger female, older female, younger male, or older male (and no other data about them).

```

predicted.data_Age_Sex <- data.frame(
  probability.of.Data=logistic_Age_Sex$fitted.values,
  AgeCategory=data$AgeCategory, Sex=data$Sex)

## We can plot the data...
ggplot(data=predicted.data_Age_Sex, aes(x=AgeCategory, y=probability.of.Data, fill=Sex)) +
  geom_point(aes(color=Sex), size=5) +
  xlab("Age") +
  ylab("Predicted probability of Death in Acute Heat in all birds") +
  theme_classic()

```



```
## Since there are only two probabilities (one for females and one for males),
## we can use a table to summarize the predicted probabilities.
xtabs(~ probability.of.Data + AgeCategory + Sex, data=predicted.data_Age_Sex)
```

```
## , , Sex = F
##
##
##      AgeCategory
## probability.of.Data  Older  Younger
## 6.93202813993879e-10      0      0
## 4.35727481485329e-09      0      0
## 0.0434782608695653     23      0
## 0.222222222222222      0     45
##
## , , Sex = M
##
##
##      AgeCategory
## probability.of.Data  Older  Younger
## 6.93202813993879e-10      5      0
## 4.35727481485329e-09      0     14
## 0.0434782608695653      0      0
## 0.222222222222222      0      0
```

Read in the data - Younger animals only

```
YoungerData=read.csv("Younger.csv")
xtabs(~ Died_InTrt + Sex, data=YoungerData)
```

```
##           Sex
## Died_InTrt F  M
##           0 36 14
##           1  9  0
```

Run a generalized linear model for sex in the dataframe only containing younger individuals

```
logistic3 <- glm(Died_InTrt ~ Sex, data=YoungerData, family="binomial")
summary(logistic3)
```

```
##
## Call:
## glm(formula = Died_InTrt ~ Sex, family = "binomial", data = YoungerData)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.66805  -0.66805  -0.66805  -0.00013   1.79412
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  -1.3863      0.3727  -3.72 0.000199 ***
## SexM         -17.1798    1743.2485  -0.01 0.992137
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 50.397  on 58  degrees of freedom
## Residual deviance: 45.036  on 57  degrees of freedom
## AIC: 49.036
##
## Number of Fisher Scoring iterations: 17
```

```
Youngfemale.log.odds <- log(9 / 36)
Youngfemale.log.odds
```

```
## [1] -1.386294
```

```
Youngmale.log.odds <- log(0 / 18)
Youngmale.log.odds
```

```
## [1] -Inf
```

Now calculate the overall “Pseudo R-squared” and its p-value for younger individuals by sex

```
## NOTE: Since we are doing logistic regression...
## Null deviance = 2*(0 - LogLikelihood(null model))
##               = -2*LogLikelihood(null model)
## Residual deviance = 2*(0 - LogLikelihood(proposed model))
##               = -2*LogLikelihood(proposed model)
ll.null <- logistic3$null.deviance/-2
ll.proposed <- logistic3$deviance/-2

## McFadden's Pseudo R^2 = [ LL(Null) - LL(Proposed) ] / LL(Null)
(ll.null - ll.proposed) / ll.null
```

```
## [1] 0.1063724
```

```
## chi-square value = 2*(LL(Proposed) - LL(Null))
## p-value = 1 - pchisq(chi-square value, df = 2-1)
1 - pchisq(2*(ll.proposed - ll.null), df=1)
```

```
## [1] 0.02059365
```


```
1 - pchisq((logistic3$null.deviance - logistic3$deviance), df=1)
```

```
## [1] 0.02059365
```

Lastly, let's see what this logistic regression predicts.

```
predicted.data <- data.frame(
  probability.of.YoungerData=logistic3$fitted.values,
  sex=YoungerData$Sex)

## We can plot the data...
ggplot(data=predicted.data, aes(x=sex, y=probability.of.YoungerData)) +
  geom_point(aes(color=sex), size=5) +
  xlab("Sex") +
  ylab("Predicted Probability of Acute Heat Death in Younger Birds") +
  theme_classic()
```



```
## Since there are only two probabilities (one for females and one for males),
## we can use a table to summarize the predicted probabilities.
xtabs(~ probability.of.YoungerData + sex, data=predicted.data)
```

```
##
##      probability.of.YoungerData      sex
##      probability.of.YoungerData  F  M
##      8.6468692450059e-09  0 14
##      0.2                  45  0
```